

Ai-Chun Pang / Instructor

林偉哲 陳令原 / T.A.s

Assignment 2 - TCP Socket Programming

Assignment 2

Announcement

Specification (1/4)

- In this assignment, you need to implement a simple Network Storage System, with these functions:
 - Client can **watch a “.mpg” video (streaming)** on the server
 - Client can **upload files** to server
 - Client can **download files** from server
 - Client can **list what files are in the folder of server**
- For video steaming, you **don't need to send audio**. You can just send frames in **RAW format**.
- After upload and download, you need to **ensure the files are identical between source and destination**.
- In this assignment, all the transmission should be implemented by **the socket of TCP**.

Specification (2/4)

- You are required to write a Makefile for compilation. Thus, the command should be

```
$ make client           // for client code
```

```
$ make server           // for serve code
```

- After compilation, there should be 2 binary files named “client” and “server”
- When we launch the server app, we will enter

```
$ ./server [port]       // [port] will be determined
```

- When we launch the client app, we will enter

```
$ ./client [ip:port]    // ip is the ip address of server  
                        port is determined by the  
                        command above
```

- After the launch, client or server is required to **create their own folders while any folder doesn't exist.**

Specification (3/4)

- Server is required to support multiple connections. That is, there **can be more than 1 client** connecting to the server **simultaneously**.
- After building up of a connection, a user can enter the commands below on the client,

\$ ls // Then, the client's will print out what files is in the server's folder.

\$ put <filename> // Then, the client will upload the file with the <filename> to the server's folder

\$ get <filename> // Then, the client will download the file with the <filename> to the client's folder

\$ play <videofile> // Then, the client will play the <videofile> from the server to the client

Specification (4/4)

- If the command doesn't exist or the command format is wrong, please print out “**Command not found.**” or “**Command format error.**” on the client.
- If the file doesn't exist while putting or getting a file, please print out “**The '<filename>' doesn't exist.**” on the client.
- If the video file is not a “.mpg” file while playing a video file, please print out “**The '<videofile>' is not a mpg file.**”
- Client should be able to send another command after a command is finished.
- The multiple connections should be implemented with **pthread.h**, while the video player should be implemented with **OpenCV**.
- The implementation must be in **C or C++**.

Grading Policy (1/3)

- This assignment accounts for 12% of the total score.
- **Command Sending** (10%)
 - The client sends commands correctly (5%)
 - The client prints out responses correctly (5%)
- **Video Streaming** (25%)
 - Correctly playing a 720p Video (1280*720) (15%)
 - Correctly playing a resolution-unknown video (10%)
(That is, client has no idea about the resolution of the video before requesting a video to play.)
- **File Transferring** (25%)
 - There would be 6 files with different sizes (5%*5)
(You will get 0 point in a testcase if the transfer of a testcase is terminated or halts before finished, or the files are not identical between source and destination after the transfer.)

Grading Policy (2/3)

- **Multiple Connections** (20%)
 - Use <pthread.h> to achieve this function (basic) 10%
 - Use select() to achieve this function (advance) 20%
 - You just need to choose one of above to implement.
- **Report** (20%)
 - Draw a flowchart of the video streaming and explains how it works in detail. (5%)
 - Draw a flowchart of the file transferring and explains how it works in detail. (5%)
 - What is SIGPIPE? It is possible to happen to your code? If so, how do you handle it? (5%)
 - Is blocking I/O equal to synchronized I/O? Please give me some examples to explain it. (5%)

Grading Policy (3/3)

- **Submission**

- Your report format must be in **“.pdf”** format and named “report.pdf”, or else **you will get zero point** in the part.
- Please put all the files **into a folder** named **hw2_<student id>**, and compress the folder as a **.zip** file, and then submit the **.zip** file to **here**. The password is **<student id>** (alphabet is in uppercase).
- If we **cannot compile or execute your code**, you will **have a chance to demo your results** in your own environment.
- The penalty for **wrong format** is **10 points**.
- **No plagiarism** is allowed. A plagiarist will be graded **zero**.

- **Deadline**

- Soft Deadline: **23:59:59, November 12th, 2019**
- Hard Deadline: **23:59:59, November 19th, 2019** (**get *0.9 points**)
- Penalty for late submission after hard deadline is **“20 points per day”** (**after *0.9**)

Environment Setup

Environment

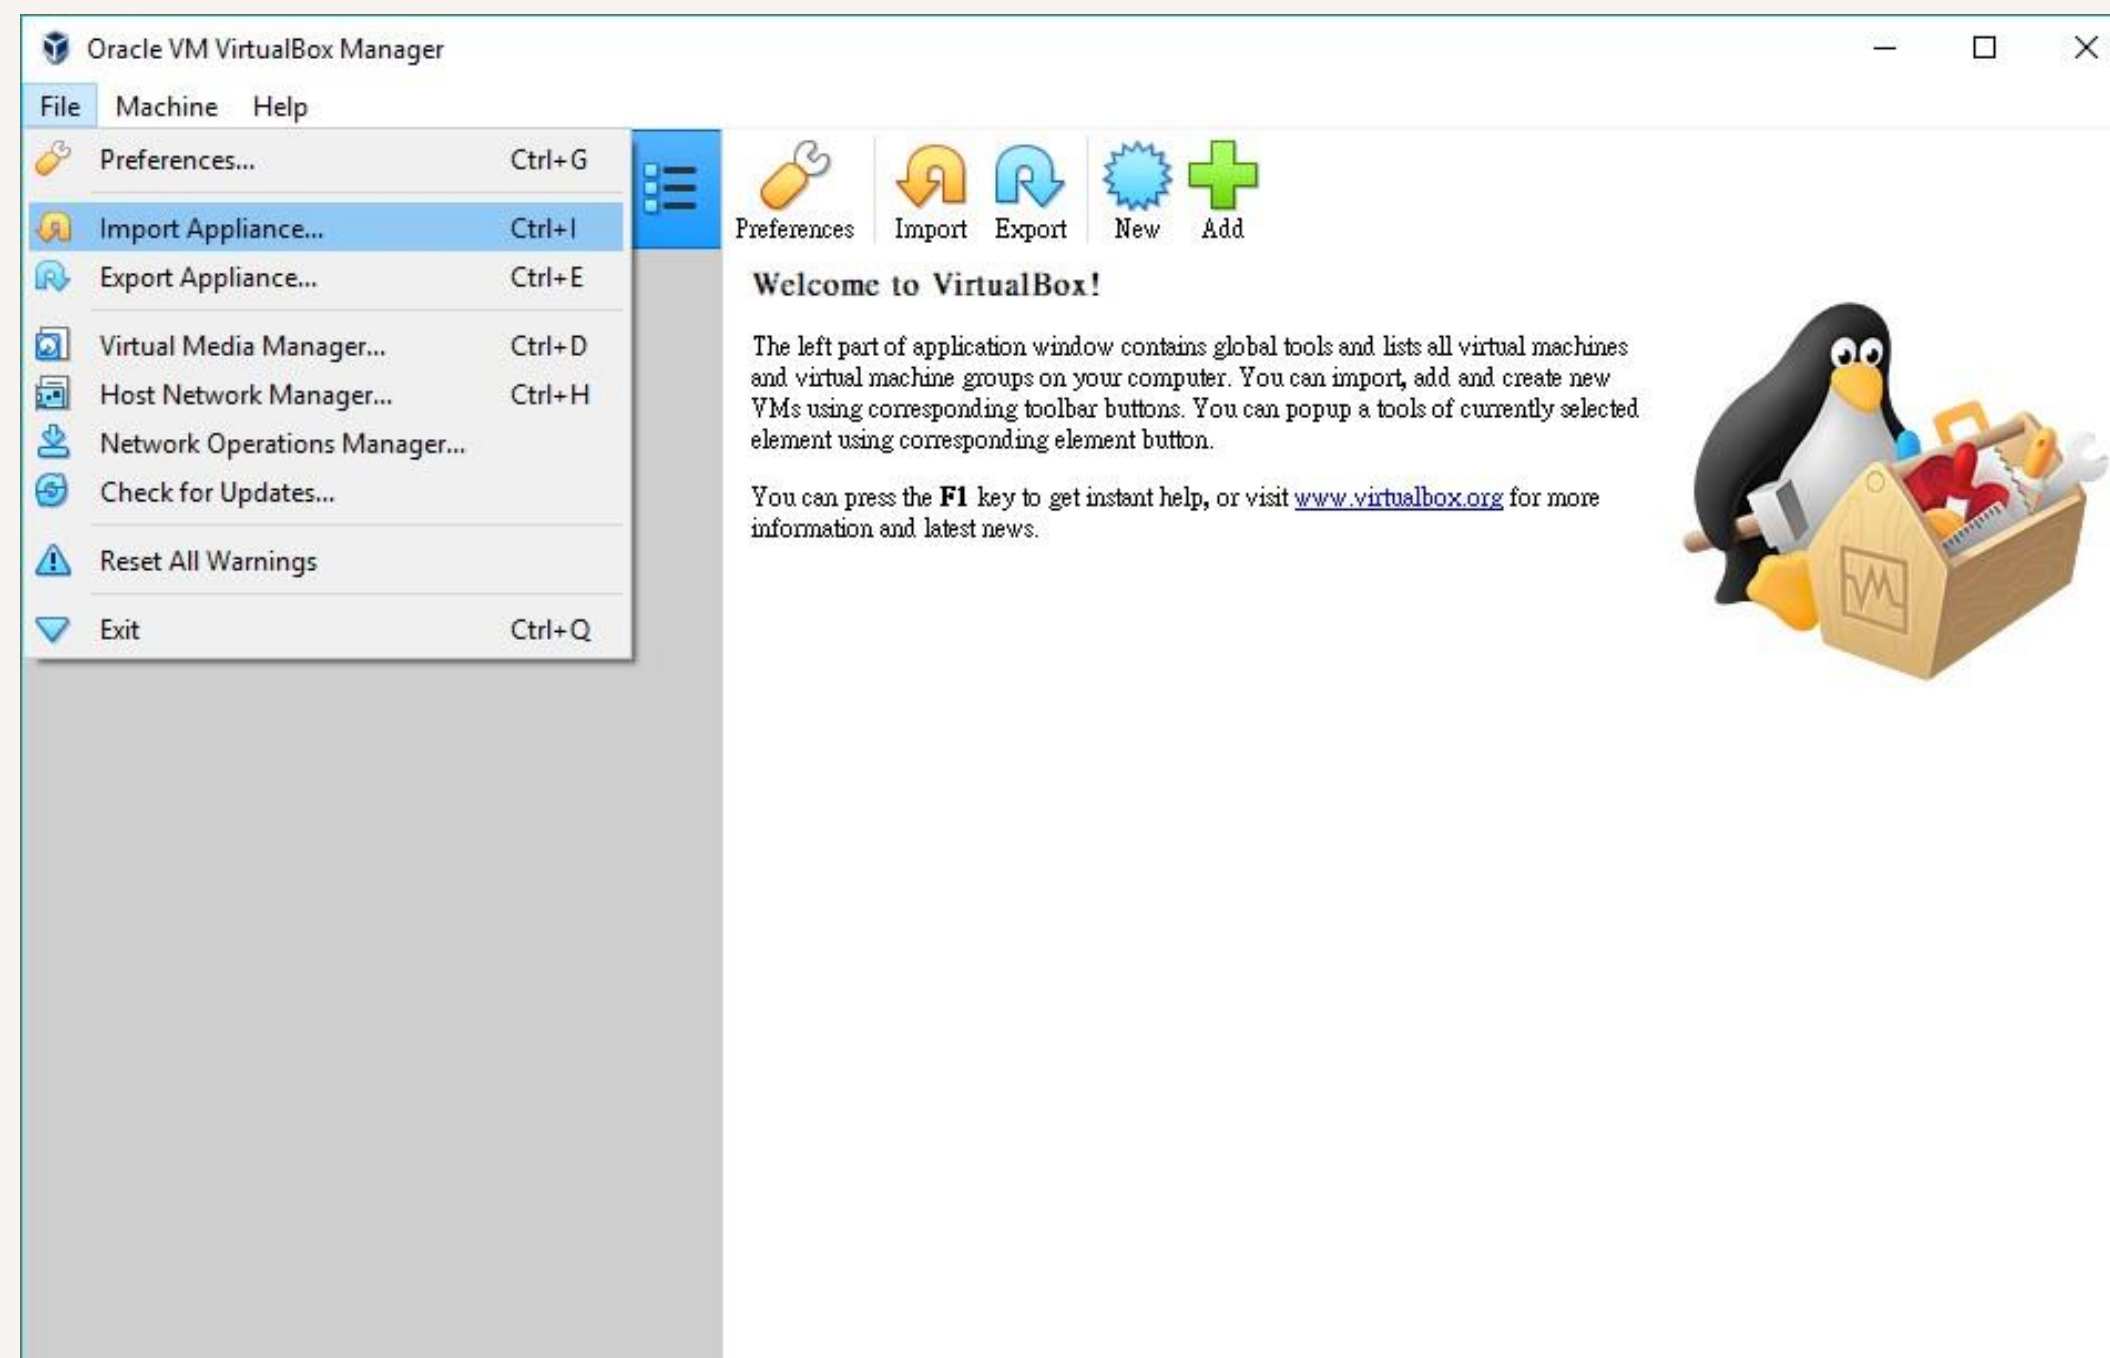
- We provide a VirtualBox environment for you to run our binary code and you can run Wireshark on this environment.
- If you would like to setup the environment on your OS rather than our virtual machine, here is information of our environment
 - Ubuntu 16.04 x64
 - OpenCV 3.3.1 (will be also required in later Assignments)
- You can install OpenCV 3.3.1 by following the instruction [here](#).

VirtualBox Setup

- Download the **VM** from
 - our server,
 - our Google Drive,
- Install Virtualbox (natively installed on the computers of Lab R204).

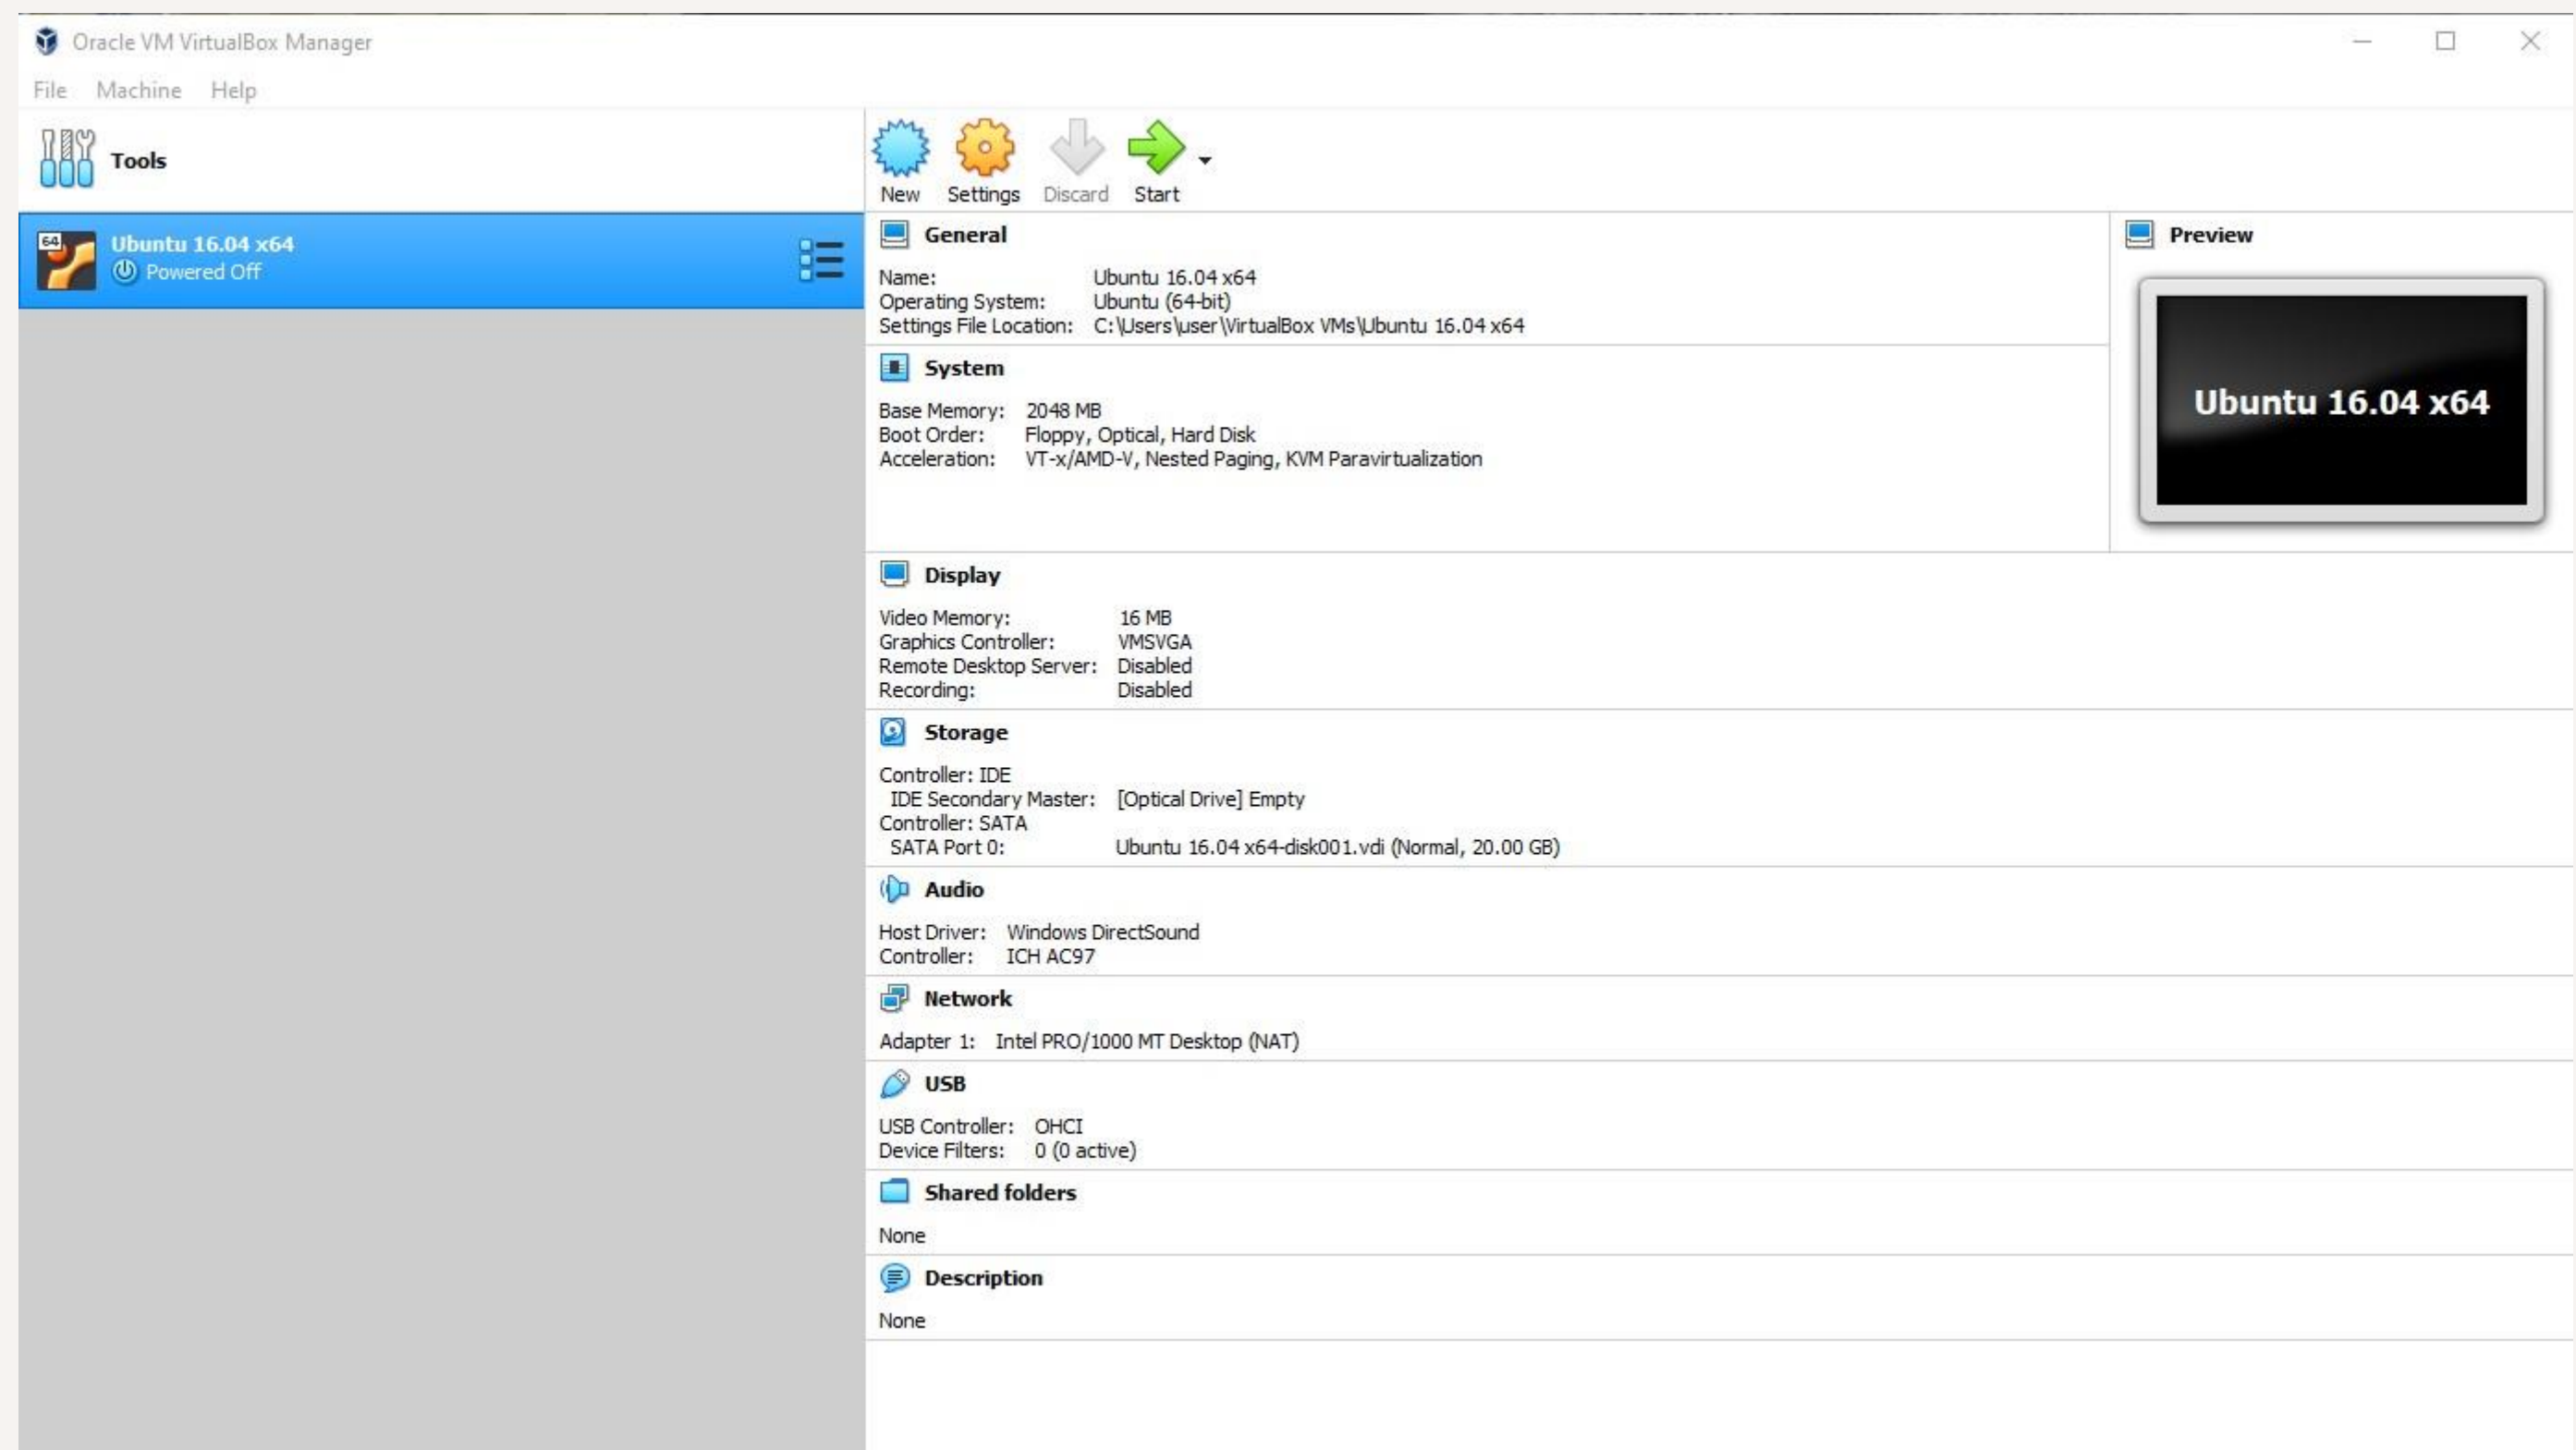
VirtualBox Setup

- Go to “File” and click “Import Appliance” to import the “CN-Ubuntu_16.04_x64.ova”



VirtualBox Setup

- Choose “Ubuntu 16.04 x64” and then start the machine.



Auxiliary Libraries

OpenCV

- Is an opensource library for computer vision.
- Mat is an image container to load an image so that you can easy to do image processing, recognition, etc.
- In this assignment, we use this library to get frames from videos on server, and show frames on client.
- An example code is [here](#). [Here](#) is an .mpg video file if you need.
- To compile code with OpenCV,

```
$ g++ <file name> -o <output name> \  
`pkg-config --cflags --libs opencv`
```


Pthread

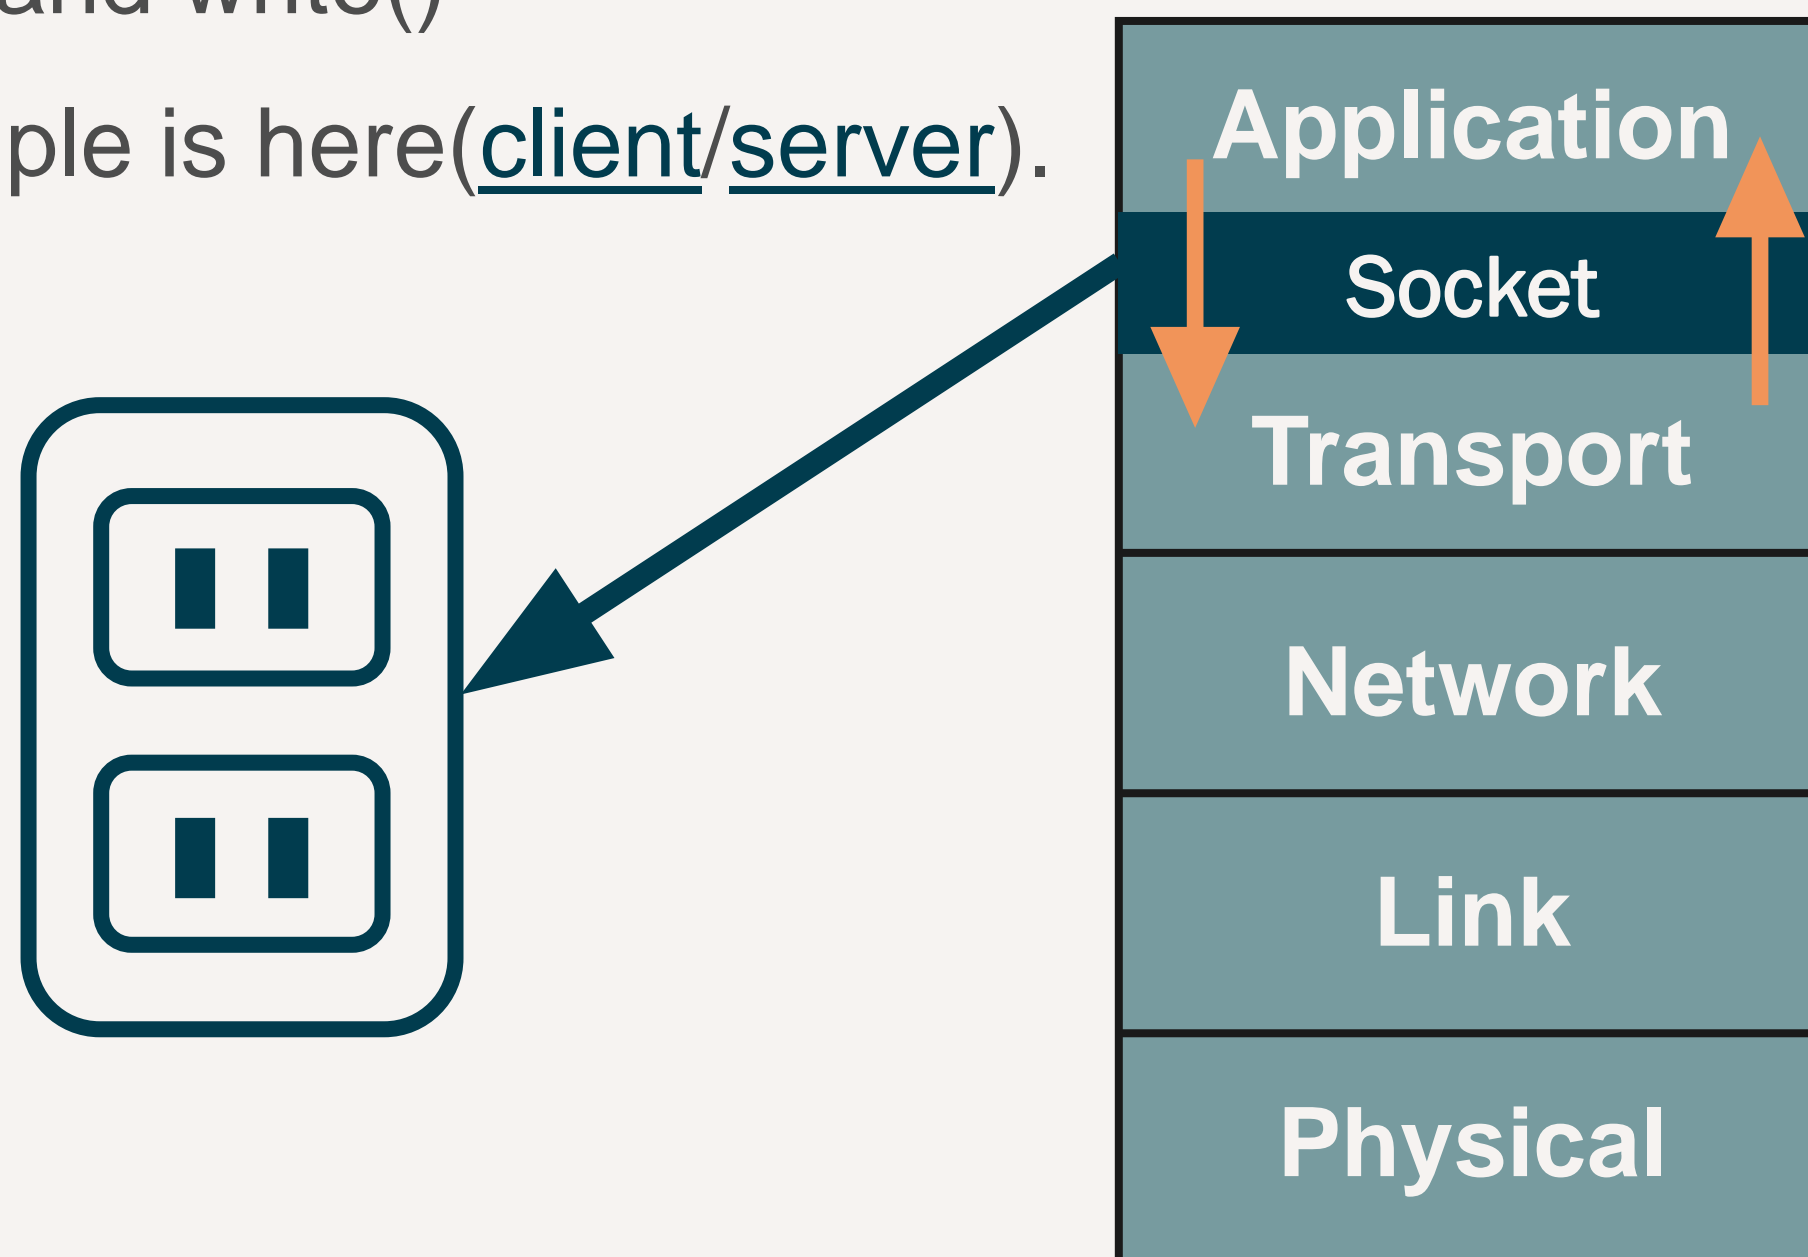
- Pthead, i.e., POSIX Thread, is used to implement multi-thread parallelization in POSIX environment.
- You can use pthread to achieve multiple connections.
- You don't need to deal with synchronization issues, i.e., in our testcases, it won't put a video with the same file name.
- An example code is [here](#).
- To compile with Pthread,

```
$ g++ <file name> -o <output name> -pthread
```

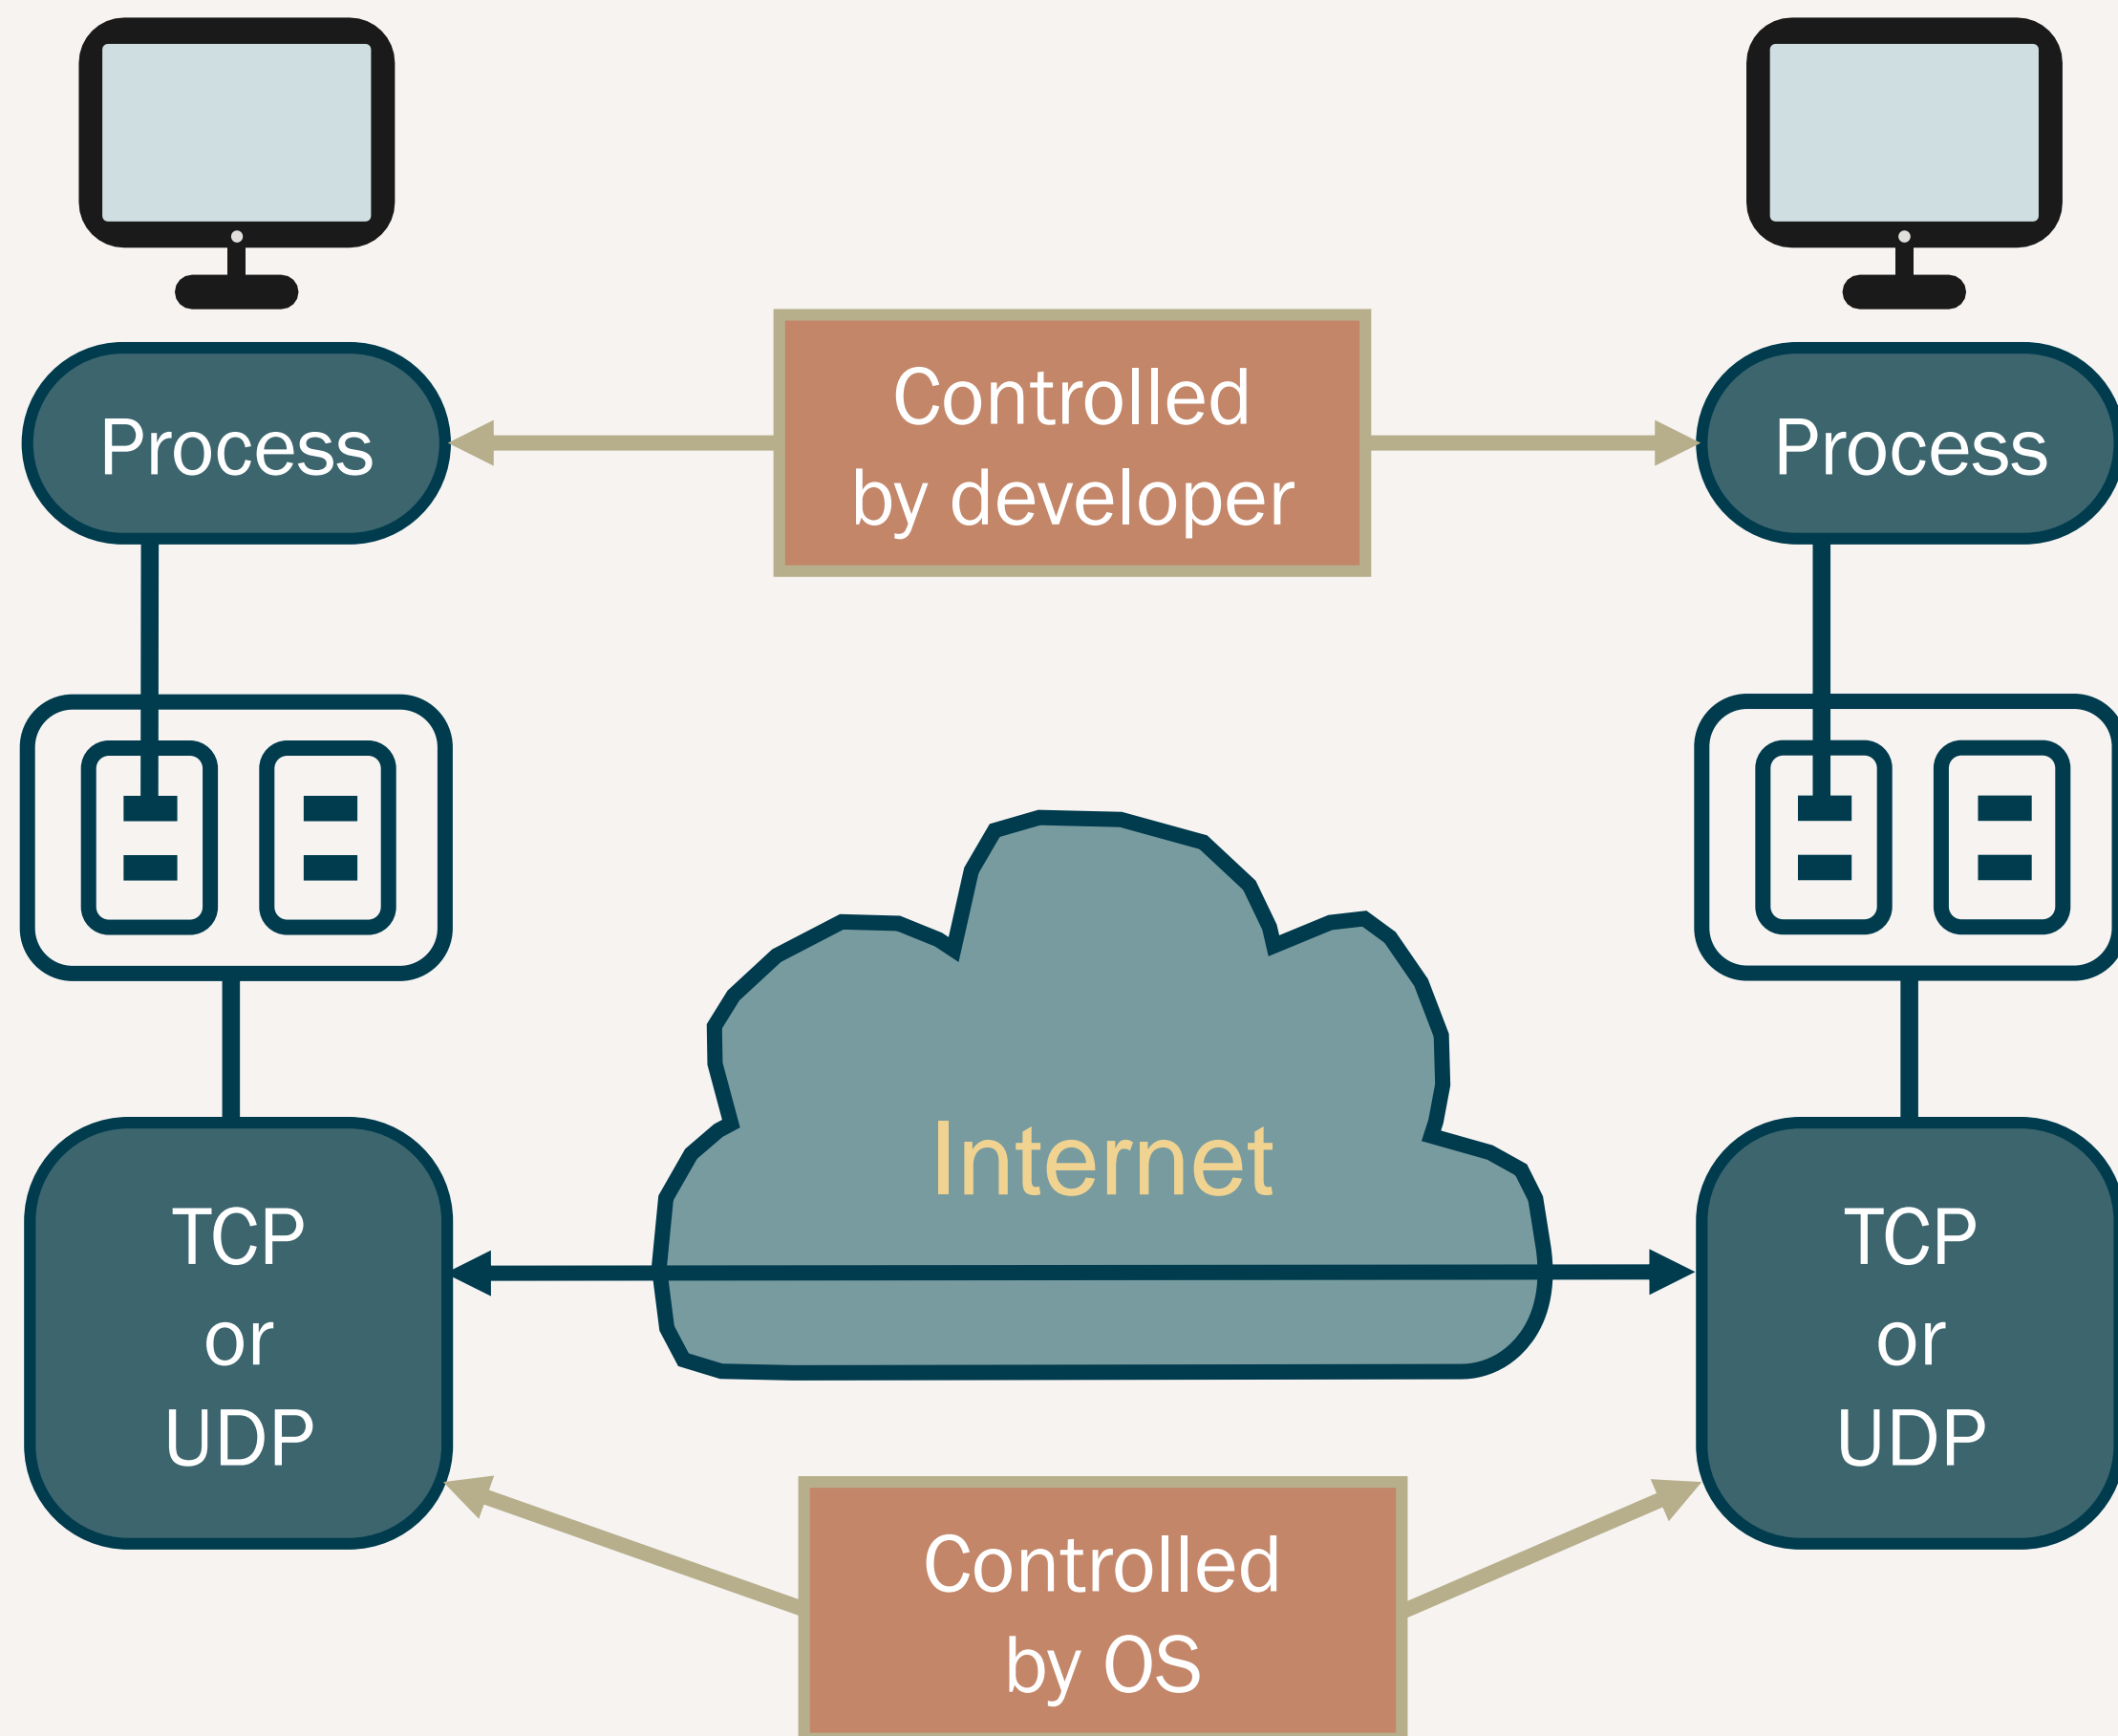
Socket Programming Tutorial

What is Socket?

- Socket is the **API for the TCP/IP protocol stack**.
- Provides communication **between the Application layer and Transport layer**
- Make internet communication **like a file descriptor**.
 - read() and write()
- An example is here([client/server](#)).



What is Socket?



File Descriptors

- When we open an existing file or create a new file, the kernel return a file descriptor to the process.
- If we want to read or write a file, we identify the file with the file descriptor.

Integer value	Name	<unistd.h> symbolic constant	<stdio.h> file stream
0	Standard input	STDIN_FILENO	stdin
1	Standard output	STDOUT_FILENO	stdout
2	Standard error	STDERR_FILENO	stderr

```
FILE *fp = fopen("this.txt", "w");
```

```
fprintf(fp, "Happy Coding.");
```

```
fclose(fp);
```

```
printf("This is Computer Networking.\n");
```


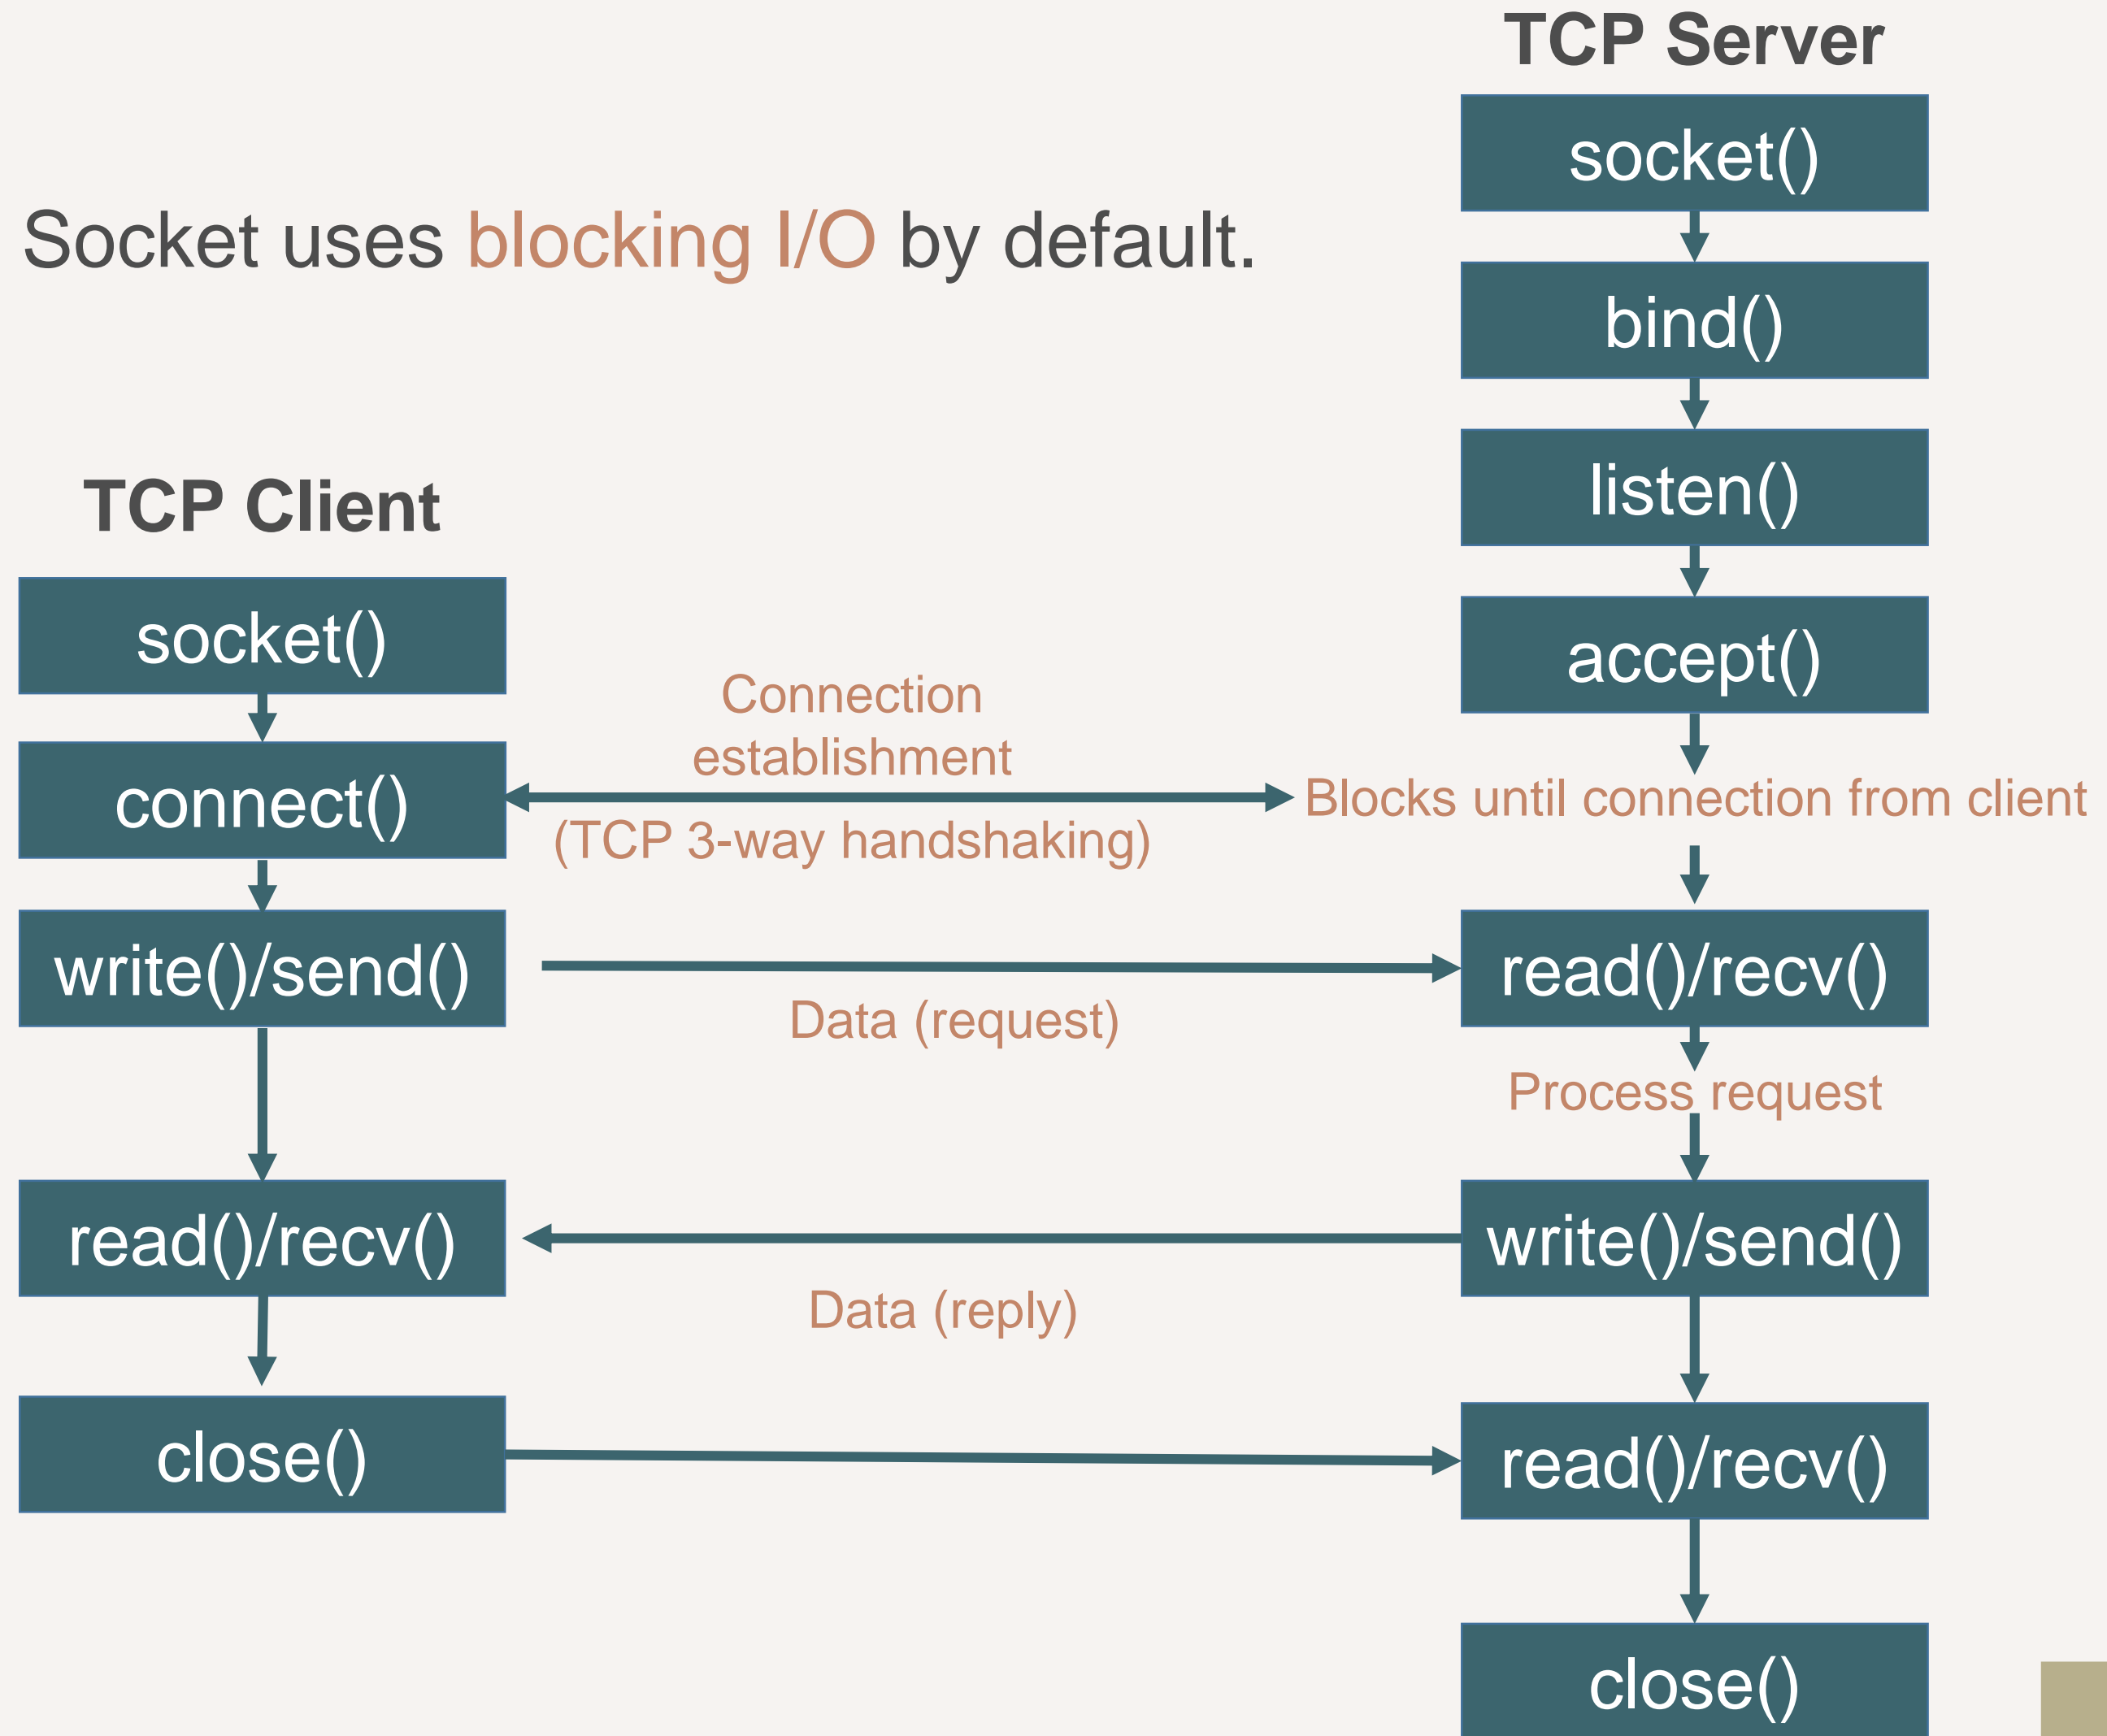
```
fprintf(stdout, " This is Computer Networking.\n");
```

TCP Service

- **TCP (Transmission Control Protocol)**
 - **Connection-oriented**
 - **Reliable** transport
 - Flow control
 - Congestion control
- What is Socket-Address?
 - **IP address + Port number**
 - IP address: To find out the machine (**Network Layer**)
 - Port number: To find out the process (Transport Layer)

TCP Socket Flow Chart

- Socket uses **blocking I/O** by default.



socket ()

- Create the endpoint for connection

```
#include <sys/types.h, sys/socket.h>
int socket (int domain, int type, int protocol);
           // returns a file descriptor if it's success; -1 otherwise
```

- domain
 - AF_UNIX/AF_LOCAL: communication between 2 processes on a host, so that they can share a file system.
 - AF_INET, AF_INET6: communication between processes on different hosts through the Internet. AF_INET is for IPv4, where AF_INET6 is for IPv6
- type
 - SOCK_STREAM: sequential and connection-oriented (TCP)
 - SOCK_DGRAM: datagram (UDP)
- protocol: defined in /etc/protocols, usually set to 0
- return: socket file descriptor (an integer)



bind()

- Bind the address to the socket

```
#include <sys/types.h, sys/socket.h>
int bind (int sockfd, struct sockaddr *addr, socklen_t len);
// returns 0 if it's success; -1 otherwise
```

- sockfd: specifies the socket file descriptor to bind.
- sockaddr
 - specifies the socket address to be associated with sockfd
 - You can use “struct sockaddr_in*” defined in <netinet/in.h>, and then cast it into “struct sockaddr”
- len: specifies the size of sockaddr (=sizeof(struct sockaddr))

```
struct sockaddr_in{
    short sin_family;           // address family. EX:AF_INET
    unsigned short sin_port;    // port number for network
    struct in_addr sin_addr;    // IP address for network
    unsigned char sin_zero[8];  // pad to sizeof(struct sockaddr)
}
```



listen()

- **Listen** for connections on a socket

```
#include <sys/types.h, sys/socket.h>
int listen (int sockfd, int backlog); // returns 0 if it's success; -1 otherwise
```

- sockfd: specifies the socket file descriptor to listen.
- backlog: specifies the number of users allowed in queue.
Linux typically add **3** to the number specified,
while other OS has different implementations.



accept()

- **Accept** the connection on a socket. After accepting the connection, it **creates a new file descriptor** for the client. The original socket is not affected.

```
#include <sys/types.h, sys/socket.h>
int accept (int sockfd, struct sockaddr *addr, socklen_t *addrlen);
// returns a file descriptor if it's success; -1 otherwise
```

- **addrlen** : pointer to the length of **sockaddr**
- **Blocking** until a user **connect()** call is received.
- Format is the same as **socket()**.



connect()

- Connect to the socket from client to server

```
#include <sys/types.h, sys/socket.h>
int connect (int sockfd, struct sockaddr *addr, socklen_t len);
// returns 0 if it's success; -1 otherwise
```

- Format is the same as bind().



close()

- Close the file descriptor

```
#include <unistd.h>  
int close (int sockfd);
```

// returns 0 if it's success; -1 otherwise



read()/recv()

- Read data from socket file descriptor

```
#include <unistd.h>
ssize_t recv (int fd, void *buf, size_t len, int flag);
//return : number of bytes read if it's success, -1 otherwise
```

- fd: specifies the socket file descriptor to read data from
- buf: specifies the buffer to contain the received data
- count: specifies the size to receive
- flag: (read()) has no this parameter.) It's about some details like blocking/nonblocking.
- Reading data from file may be
 - Success
 - EOF (end of file) (i.e., return = 0)
- It may be blocked. (block I/O)



write()/send()

- Write data to socket file descriptor

```
#include <unistd.h>
ssize_t send (int fd, void *buf, size_t len, int flag);
//return : number of bytes written if it's success, -1 otherwise
```

- fd: specifies the socket file descriptor to send data to
- buf: specifies the buffer to contain the data to transmit
- count: specifies the size to send
- flag: (write() has no this parameter.) It's about some details.
- Reading data from file may be
 - Success
 - EOF (end of file) (i.e., return = 0)
 - Error (i.e., return < 0)
- It may be blocked. (block I/O)



Useful Functions

- Address and port numbers are stored as integers.
Different machines implements **different endian**.
They may communicate with each other on the network.
- IP address is usually hard to remember.
We need to **translate hostname to IP address**.

Useful Functions

- Converting IP address and port number
 - htonl() : for IP address (host -> network)
 - ntohl() : for IP address (network -> host)
 - htons() : for port number (host -> network)
 - ntohs() : for port number (network -> host)
- Translate a hostname to IP address

```
#include <netdb.h>
```

```
struct hostent *gethostbyname (const char *name);
```

```
// return: host environment if it's success, NULL otherwise
```

Supplementary Materials

How to Simulate a Bad Network

- Some bugs may occur when the network is not good.
- In case of bad network, we should simulate a bad network to test our program in advance.

- Linux

- Get the network interfaces list in your machine

```
$ ifconfig
```

- Then, to make your personal internet slow, you can enter

```
$ sudo tc qdisc add dev <interface> root netem delay 500ms
```

In this way, the delay will be 500ms.

- You can turn it off with

```
$ sudo tc qdisc del dev <interface> root netem
```

How to Trace Kernel

- Sometimes, the service may terminate without any error message.
- It happens usually because of some kernel issues.
- To trace the interactions between your code and kernel, we can use **strace**.
- To run your program with strace, you can enter

```
$ strace ./<name>
```

Behavior of `send()` and `recv()`

- In fact, a `send()` doesn't imply all the data in the buffer are sent.
- In addition, a `send()` doesn't imply all the data in the buffer are sent in a packet, and even 2 `send()` don't imply they are in different packets.
- You are required to design a protocol so that each receive has the same size as the send in respect to it.

select()

- Select() provides you to supervise multiple sockets, telling you which is able to read or write, etc.
- With select(), it is possible to achieve Asynchronous Blocking I/O.
- For more details about I/O, you can go to [this website](#).
[\(Report 4\)](#)
- If you want to implement this assignment with select(), please refer to [this website](#).

select()

- **Monitor** whether there is at least one fd available

```
#include <unistd.h>
int select(int nfd, fd_set*, readfds, fd_set* writefds, fd_set* exceptfds,
struct timeval* timeout); //return : 1 if it's success, -1 otherwise
```

- nfd: specifies number of file descriptors to monitor.
- readfds: specifies the pointer to read file descriptor list
- writefds: specifies the pointer to write file descriptor list
- exceptfds: specifies the pointer to error file descriptor list
- timeout: deadline for select().

select()

```
void FD_SET(int fd, fd_set *set);  
void FD_CLR(int fd, fd_set *set);  
int FD_ISSET(int fd, fd_set *set); // return: 1 if it's available, else: 0  
void FD_ZERO(fd_set *set);
```

- FD_SET: Add the file descriptor into the set
- FD_CLR: Remove the file descriptor from the set
- FD_ISSET: Check if the file descriptor is available
- FD_ZERO: Clear the set

Reference

- [Beej's Guide to Network Programming \(中文\)](#)
- [Beej's Guide to Network Programming \(English\)](#)